

```
In [ ]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

Loading Data

```
In [ ]: 1 df = pd.read_csv('Seasonal Data.csv', encoding="ISO-8859-1")
        2 df.head()
```

```
Out[15]:
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	Customer Name	Segment	Country	City	...	Postal Code	Region	P
0	1	CA-2016-152156	11/8/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	42420	South	FL 101
1	2	CA-2016-152156	11/8/2016	11/11/2016	Second Class	CG-12520	Claire Gute	Consumer	United States	Henderson	...	42420	South	FL 101
2	3	CA-2016-138688	6/12/2016	6/16/2016	Second Class	DV-13045	Darrin Van Huff	Corporate	United States	Los Angeles	...	90036	West	CA 101
3	4	US-2015-108966	10/11/2015	10/18/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	...	33311	South	FL 101
4	5	US-2015-108966	10/11/2015	10/18/2015	Standard Class	SO-20335	Sean O'Donnell	Consumer	United States	Fort Lauderdale	...	33311	South	CA 101

5 rows × 21 columns

Data Cleaning

```
In [ ]: 1 df = df.drop_duplicates()
        2
        3 df['Order Date'] = pd.to_datetime(df['Order Date'])
        4
        5 print(df.isnull().sum())
```

```
Row ID      0
Order ID    0
Order Date  0
Ship Date   0
Ship Mode   0
Customer ID 0
Customer Name 0
Segment     0
Country     0
City        0
State       0
Postal Code 0
Region      0
Product ID  0
Category    0
Sub-Category 0
Product Name 0
Sales       0
Quantity    0
Discount    0
Profit      0
dtype: int64
```

```
In [ ]: 1 monthly_data = df[['Order Date', 'Profit']].copy()
2 monthly_data.set_index('Order Date', inplace=True)
3
4 # Resample to monthly frequency and mean profit
5 monthly_profit = monthly_data.resample('ME').mean()
6 print(monthly_profit)
```

	Profit
Order Date	
2014-01-31	31.015072
2014-02-28	18.745835
2014-03-31	3.176624
2014-04-30	25.843224
2014-05-31	22.448439
2014-06-30	36.863144
2014-07-31	-5.884494
2014-08-31	34.758856
2014-09-30	31.074998
2014-10-31	21.687153
2014-11-30	29.220525
2014-12-31	32.315000
2015-01-31	-56.569086
2015-02-28	43.966419
2015-03-31	70.522448
2015-04-30	26.171851
2015-05-31	31.971705
2015-06-30	24.170704
2015-07-31	23.490345
2015-08-31	33.684330
2015-09-30	28.017620
2015-10-31	16.972084
2015-11-30	38.502433
2015-12-31	25.370145
2016-01-31	31.739588
2016-02-29	60.296139
2016-03-31	22.159313
2016-04-30	17.516558
2016-05-31	38.498428
2016-06-30	23.871247
2016-07-31	22.054119
2016-08-31	11.716303
2016-09-30	25.698781
2016-10-31	82.873176
2016-11-30	10.841642
2016-12-31	50.810538
2017-01-31	46.067349
2017-02-28	15.082916
2017-03-31	61.982737
2017-04-30	4.597488
2017-05-31	26.209020
2017-06-30	33.564636
2017-07-31	30.763811
2017-08-31	41.472274
2017-09-30	23.946744
2017-10-31	31.125086
2017-11-30	21.111337
2017-12-31	18.362223


```
In [ ]: 1 weekly_data = df[['Order Date', 'Profit']].copy()
2 weekly_data.set_index('Order Date', inplace=True)
3
4 #Covertng into weekly profit
5 weekly_profit = weekly_data.resample('W').mean()
6 print(weekly_profit)
```

```
Profit
Order Date
2014-01-05 -11.110980
2014-01-12  82.671462
2014-01-19  18.131195
2014-01-26  26.139231
2014-02-02  21.776973
...
2017-12-03  16.013970
2017-12-10   8.390592
2017-12-17   9.856660
2017-12-24  42.883730
2017-12-31  20.707991
```

[209 rows x 1 columns]

Plot

```
In [ ]: 1 plt.figure(figsize=(14, 6))
2 sns.lineplot(data=weekly_profit, x=weekly_profit.index, y='Profit')
3 plt.title('Weekly Profit Over Time')
4 plt.xlabel('Date')
5 plt.ylabel('Weekly Profit')
6 plt.grid(True)
7 plt.tight_layout()
8 plt.show()
9
```



```
In [ ]: 1 from statsmodels.tsa.seasonal import seasonal_decompose
2
3 # Decompose the time series
4 decomposition = seasonal_decompose(weekly_profit, model='additive')
5 decomposition.plot()
6 plt.figure(figsize=(12, 8))
7 plt.tight_layout()
8 plt.suptitle('Seasonal Decomposition of Weekly profit')
9 plt.show()
10
```



<Figure size 1200x800 with 0 Axes>

From the graph we can see that this data is seasonal dataset


ADF test

```
In [23]: 1 from statsmodels.tsa.stattools import adfuller
2
3 result = adfuller(weekly_profit['Profit'])
4
5 print("ADF Statistic:", result[0])
6 print("p-value:", result[1])
7 for key, value in result[4].items():
8     print(f'Critical Value ({key}): {value}')
9
10 if result[1] < 0.05:
11     print("The series is stationary.")
12 else:
13     print("The series is not stationary.")
14
```

```
ADF Statistic: -13.790688189145076
p-value: 8.938098275917793e-26
Critical Value (1%): -3.4621857592784546
Critical Value (5%): -2.875537986778846
Critical Value (10%): -2.574231080806213
The series is stationary.
```


```
In [24]: 1 from statsmodels.graphics.tsaplots import plot_acf
2
3 # ACF
4 plt.figure(figsize=(12, 6))
5 acf_plot_close = plot_acf(weekly_profit.dropna(), lags=48, alpha=0.05) # confidence inter
6
7 # Add a dotted line at the significance threshold
8 plt.axhline(y=-1.96/np.sqrt(len(weekly_profit)), linestyle='--', color='gray')
9 plt.axhline(y=1.96/np.sqrt(len(weekly_profit)), linestyle='--', color='gray')
10
11 plt.title('Autocorrelation Function (ACF) Plot for Weekly Profit')
12 plt.show()
```

<Figure size 1200x600 with 0 Axes>



```
In [25]: 1 from statsmodels.graphics.tsaplots import plot_pacf
2
3 #PACF
4 plt.figure(figsize=(12, 6))
5 pacf_plot_temp = plot_pacf(weekly_profit.dropna(), lags=48, alpha=0.05)
6
7 # Add a dotted line at the significance threshold
8 plt.axhline(y=-1.96/np.sqrt(len(weekly_profit)), linestyle='--', color='gray')
9 plt.axhline(y=1.96/np.sqrt(len(weekly_profit)), linestyle='--', color='gray')
10
11 plt.title('Partial Autocorrelation Function (PACF) Plot for Weekly Profit')
12 plt.show()
```

<Figure size 1200x600 with 0 Axes>



```
In [26]: 1 import statsmodels.api as sm
2 import itertools
3 from statsmodels.tsa.statespace.sarimax import SARIMAX
4
5 p=d=q= range(0,2)
6 # Generate all different combinatins of p,d and q triplets
7 pdq= list(itertools.product(p,d,q))
8
9 seasonal_pdq = [(x[0], x[1], x[2],12) for x in list(itertools.product(p,d,q))]
10
11 for param in pdq:
12     for param_seasonal in seasonal_pdq:
13         try:
14             mod= sm.tsa.statespace.SARIMAX(weekly_profit,
15                                             order= param,
16                                             seasonal_order= param_seasonal,
17                                             enforce_stationarity= False,
18                                             enforce_invertibility= False)
19             results= mod.fit()
20             print(f'SARIMA{param}x{param_seasonal}12- AIC: {results.aic}')
21         except:
22             continue
```



```
SARIMA(0, 0, 0)x(0, 0, 0, 12)12- AIC: 2156.888558218995
SARIMA(0, 0, 0)x(0, 0, 1, 12)12- AIC: 2023.6058502594185
SARIMA(0, 0, 0)x(0, 1, 0, 12)12- AIC: 2082.773519163945
SARIMA(0, 0, 0)x(0, 1, 1, 12)12- AIC: 1831.74461055686
SARIMA(0, 0, 0)x(1, 0, 0, 12)12- AIC: 2020.414576116892
SARIMA(0, 0, 0)x(1, 0, 1, 12)12- AIC: 1956.241045990632
SARIMA(0, 0, 0)x(1, 1, 0, 12)12- AIC: 1902.7050713545368
SARIMA(0, 0, 0)x(1, 1, 1, 12)12- AIC: 1832.8439613077053
SARIMA(0, 0, 1)x(0, 0, 0, 12)12- AIC: 2116.09439095827
SARIMA(0, 0, 1)x(0, 0, 1, 12)12- AIC: 1999.3926962331466
SARIMA(0, 0, 1)x(0, 1, 0, 12)12- AIC: 2073.628199570655
SARIMA(0, 0, 1)x(0, 1, 1, 12)12- AIC: 1824.027617179786
SARIMA(0, 0, 1)x(1, 0, 0, 12)12- AIC: 2014.0152562383623
SARIMA(0, 0, 1)x(1, 0, 1, 12)12- AIC: 1947.2197180130447
SARIMA(0, 0, 1)x(1, 1, 0, 12)12- AIC: 1904.6318308299383
SARIMA(0, 0, 1)x(1, 1, 1, 12)12- AIC: 1825.2991589565547
SARIMA(0, 1, 0)x(0, 0, 0, 12)12- AIC: 2162.826826167911
SARIMA(0, 1, 0)x(0, 0, 1, 12)12- AIC: 2040.7427160622137
SARIMA(0, 1, 0)x(0, 1, 0, 12)12- AIC: 2214.0270070590245
SARIMA(0, 1, 0)x(0, 1, 1, 12)12- AIC: 1948.530725560736
SARIMA(0, 1, 0)x(1, 0, 0, 12)12- AIC: 2050.346362077836
SARIMA(0, 1, 0)x(1, 0, 1, 12)12- AIC: 2042.250745279631
SARIMA(0, 1, 0)x(1, 1, 0, 12)12- AIC: 2023.860774296481
SARIMA(0, 1, 0)x(1, 1, 1, 12)12- AIC: 1948.3401034852745
SARIMA(0, 1, 1)x(0, 0, 0, 12)12- AIC: 2027.8506774817458
SARIMA(0, 1, 1)x(0, 0, 1, 12)12- AIC: 1916.9678966805159
SARIMA(0, 1, 1)x(0, 1, 0, 12)12- AIC: 2068.4491286542234
SARIMA(0, 1, 1)x(0, 1, 1, 12)12- AIC: 1818.8481096506332
SARIMA(0, 1, 1)x(1, 0, 0, 12)12- AIC: 1935.3719500244456
SARIMA(0, 1, 1)x(1, 0, 1, 12)12- AIC: 1918.090667048636
SARIMA(0, 1, 1)x(1, 1, 0, 12)12- AIC: 1900.6347254313573
SARIMA(0, 1, 1)x(1, 1, 1, 12)12- AIC: 1820.7979412546983
SARIMA(1, 0, 0)x(0, 0, 0, 12)12- AIC: 2111.8269457836213
SARIMA(1, 0, 0)x(0, 0, 1, 12)12- AIC: 1996.5681732446599
SARIMA(1, 0, 0)x(0, 1, 0, 12)12- AIC: 2084.5290075328003
SARIMA(1, 0, 0)x(0, 1, 1, 12)12- AIC: 1833.7425419029564
SARIMA(1, 0, 0)x(1, 0, 0, 12)12- AIC: 1996.4727577019762
SARIMA(1, 0, 0)x(1, 0, 1, 12)12- AIC: 1958.2406970346601
SARIMA(1, 0, 0)x(1, 1, 0, 12)12- AIC: 1893.9546472056595
SARIMA(1, 0, 0)x(1, 1, 1, 12)12- AIC: 1834.8431386212847
SARIMA(1, 0, 1)x(0, 0, 0, 12)12- AIC: 2044.7281486556337
SARIMA(1, 0, 1)x(0, 0, 1, 12)12- AIC: 1927.629601302803
SARIMA(1, 0, 1)x(0, 1, 0, 12)12- AIC: 2067.723159881378
SARIMA(1, 0, 1)x(0, 1, 1, 12)12- AIC: 1827.7507321199432
SARIMA(1, 0, 1)x(1, 0, 0, 12)12- AIC: 1936.2279471303855
SARIMA(1, 0, 1)x(1, 0, 1, 12)12- AIC: 1948.9929445060156
SARIMA(1, 0, 1)x(1, 1, 0, 12)12- AIC: 1892.7456574364264
SARIMA(1, 0, 1)x(1, 1, 1, 12)12- AIC: 1826.3679193106734
SARIMA(1, 1, 0)x(0, 0, 0, 12)12- AIC: 2110.561223941431
SARIMA(1, 1, 0)x(0, 0, 1, 12)12- AIC: 1993.8679026671948
SARIMA(1, 1, 0)x(0, 1, 0, 12)12- AIC: 2163.7330649773203
SARIMA(1, 1, 0)x(0, 1, 1, 12)12- AIC: 1898.838509817877
SARIMA(1, 1, 0)x(1, 0, 0, 12)12- AIC: 1993.0632123291407
SARIMA(1, 1, 0)x(1, 0, 1, 12)12- AIC: 1994.9138377783102
SARIMA(1, 1, 0)x(1, 1, 0, 12)12- AIC: 1959.350151846037
SARIMA(1, 1, 0)x(1, 1, 1, 12)12- AIC: 1899.8898627935496
SARIMA(1, 1, 1)x(0, 0, 0, 12)12- AIC: 2029.2632925014368
SARIMA(1, 1, 1)x(0, 0, 1, 12)12- AIC: 1918.3422950270237
SARIMA(1, 1, 1)x(0, 1, 0, 12)12- AIC: 2070.2848094978062
SARIMA(1, 1, 1)x(0, 1, 1, 12)12- AIC: 1820.8416215680381
SARIMA(1, 1, 1)x(1, 0, 0, 12)12- AIC: 1926.7033175519075
SARIMA(1, 1, 1)x(1, 0, 1, 12)12- AIC: 1919.4519746027413
SARIMA(1, 1, 1)x(1, 1, 0, 12)12- AIC: 1890.0128082223362
SARIMA(1, 1, 1)x(1, 1, 1, 12)12- AIC: 1822.792107378268
```

```
In [27]: 1 Best_model= sm.tsa.statespace.SARIMAX(weekly_profit,
2                                     order= (1,0,1),
3                                     seasonal_order= (0,1,1,12),
4                                     enforce_stationarity= False,
5                                     enforce_invertibility= False)
6 results= Best_model.fit()
7 print(results.summary().tables[1])
```



```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          -0.9491      0.042     -22.667      0.000      -1.031      -0.867
ma.L1           0.9877      0.124       7.965      0.000       0.745       1.231
ma.S.L12        -1.0000      0.110     -9.120      0.000      -1.215      -0.785
sigma2         1033.9243      0.000    9.75e+06      0.000    1033.924    1033.924
=====
```

```
In [28]: 1 Best_model= sm.tsa.statespace.SARIMAX(weekly_profit,
2                                     order= (0,1,1),
3                                     seasonal_order= (0,1,1,12),
4                                     enforce_stationarity= False,
5                                     enforce_invertibility= False)
6 results= Best_model.fit()
7 print(results.summary().tables[1])
```

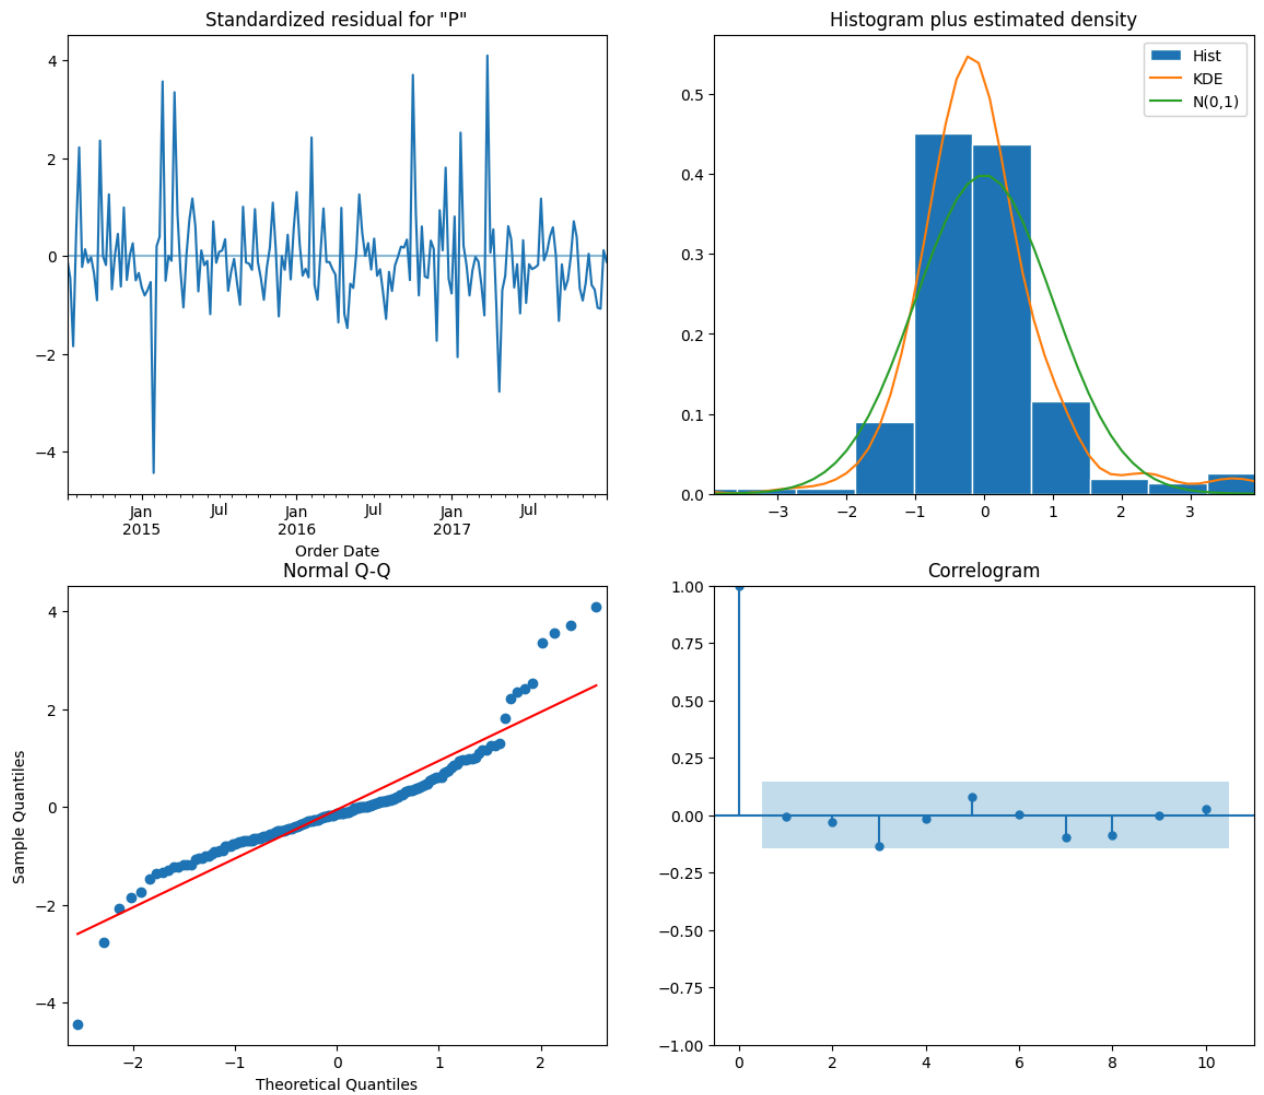
```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ma.L1          -1.0000     223.237      -0.004      0.996     -438.536      436.536
ma.S.L12        -1.0000     223.237      -0.004      0.996     -438.536      436.536
sigma2         1043.1985      0.208    5004.156      0.000    1042.790    1043.607
=====
```

```
In [29]: 1 residuals = results.resid  
2 plot_acf(residuals, lags=48)
```

Out[29]:




```
In [30]: 1 results.plot_diagnostics(figsize=(14,12))
2 plt.show()
```



```

In [31]: 1 from statsmodels.tsa.arima.model import ARIMA
2
3 # Forecast with confidence intervals
4 forecast = results.get_forecast(steps=80)
5 forecast_ci = forecast.conf_int()
6
7 # Forecast index
8 forecast_index = pd.date_range(start=weekly_profit.index[-1], periods=80, freq=weekly_prof
9
10 # Actual vs Predicted with Confidence Intervals plot
11 fig, ax = plt.subplots(figsize=(10, 6))
12 ax.plot(weekly_profit, label='Actual', color='blue')
13 ax.plot(forecast_index, forecast.predicted_mean, label='Predicted', color='red')
14
15 ax.fill_between(forecast_index,
16                 forecast_ci.iloc[:, 0],
17                 forecast_ci.iloc[:, 1], color='red', alpha=0.2, label='95% Confidence Inte
18
19 ax.legend()
20 plt.show()
21
22 # Print the summary
23 print(results.summary())

```



SARIMAX Results

=====						
Dep. Variable:	Profit		No. Observations:		209	
Model:	SARIMAX(0, 1, 1)x(0, 1, 1, 12)		Log Likelihood		-906.424	
Date:	Mon, 28 Apr 2025		AIC		1818.848	
Time:	21:45:43		BIC		1828.460	
Sample:	01-05-2014		HQIC		1822.745	
	- 12-31-2017					
Covariance Type:	opg					
=====						
	coef	std err	z	P> z	[0.025	0.975]

ma.L1	-1.0000	223.237	-0.004	0.996	-438.536	436.536
ma.S.L12	-1.0000	223.237	-0.004	0.996	-438.536	436.536
sigma2	1043.1985	0.208	5004.156	0.000	1042.790	1043.607
=====						
Ljung-Box (L1) (Q):	0.01	Jarque-Bera (JB):	222.96			
Prob(Q):	0.92	Prob(JB):	0.00			
Heteroskedasticity (H):	0.82	Skew:	0.82			
Prob(H) (two-sided):	0.45	Kurtosis:	8.17			
=====						

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).
 [2] Covariance matrix is singular or near-singular, with condition number 1.72e+22. Standard errors may be unstable.

```
In [32]: 1 df2= weekly_profit.copy()
          2
          3 df2.describe()
```

```
Out[32]:
```

	Profit
count	209.000000
mean	27.913498
std	32.653447
min	-135.678708
25%	12.506919
50%	24.582379
75%	41.007081
max	169.309043

```
In [33]: 1 from sklearn.metrics import mean_squared_error
          2
          3 n = int(len(df2) * 0.8)
          4 train = weekly_profit[:n]
          5 test = weekly_profit[n:]
          6
          7
          8 print(len(train))
          9 print(len(test))
```

```
167
42
```

```
In [34]: 1 import warnings
2 warnings.filterwarnings("ignore")
3
4 model= sm.tsa.statespace.SARIMAX(train,
5                                 order= (1,0,1),
6                                 seasonal_order= (0,1,1,12),
7                                 enforce_stationarity= False,
8                                 enforce_invertibility= False)
9 result= model.fit()
10
11 print(result.summary())
12
```

SARIMAX Results

```
=====
Dep. Variable:          Profit      No. Observations:          167
Model:                SARIMAX(1, 0, 1)x(0, 1, 1, 12)      Log Likelihood          -702.744
Date:                  Mon, 28 Apr 2025      AIC                  1413.488
Time:                  21:46:57      BIC                  1425.283
Sample:                01-05-2014      HQIC                 1418.282
                    - 03-12-2017
Covariance Type:                opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.9616      0.061     15.893      0.000      0.843      1.080
ma.L1         -1.0000     364.549     -0.003      0.998     -715.503      713.503
ma.S.L12       -1.0000     364.583     -0.003      0.998     -715.569      713.569
sigma2        1019.3438      0.266    3839.267      0.000     1018.823     1019.864
=====
Ljung-Box (L1) (Q):                0.10      Jarque-Bera (JB):          149.46
Prob(Q):                          0.75      Prob(JB):                0.00
Heteroskedasticity (H):            0.67      Skew:                    0.45
Prob(H) (two-sided):              0.17      Kurtosis:                7.96
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).
 [2] Covariance matrix is singular or near-singular, with condition number 1.53e+22. Standard errors may be unstable.

```
In [35]: 1 # Forecast using the trained model
2 forecast_result = result.forecast(steps=len(test))
3
4 mse = mean_squared_error(test, forecast_result)
5 rmse = mse**0.5
6
7 print(f'Mean Squared Error: {mse}')
8 print(f'Root Mean Squared Error: {rmse}')
9
```

Mean Squared Error: 1056.9852814141088
 Root Mean Squared Error: 32.5113100537968


```
In [36]: 1 # For cheking the model is good or not
2 test.mean(), np.sqrt(test.var())
```

```
Out[36]: (Profit      29.619405
dtype: float64,
Profit      31.46318
dtype: float64)
```

```



In [37]: 1 # Forecast with confidence intervals
2 forecast = results.get_forecast(steps=80)
3 forecast_ci = forecast.conf_int()
4
5 # Plot Train, Test, and Forecast with Confidence Intervals
6 plt.figure(figsize=(10, 6))
7 plt.plot(train, label='Train', color='blue')
8 plt.plot(test, label='Test', color='green')
9 plt.plot(forecast_result, label='ForecastSARIMA', color='red')
10
11 plt.fill_between(forecast_ci.index,
12                 forecast_ci.iloc[:, 0],
13                 forecast_ci.iloc[:, 1], color='red', alpha=0.2, label='95% Confidence Int
14
15 plt.legend()
16 plt.show()

```





```
In [ ]: 1 train.plot(legend= True, label= 'Train', figsize= (10,6))
        2 test.plot(legend= True, label= 'Test')
        3 forecast_result.plot(legend= True, label= 'ForecastSARIMA')
```

Out[29]: <Axes: xlabel='Order Date'>



```
In [38]: 1 final_model= sm.tsa.statespace.SARIMAX(weekly_profit,
2                                     order= (1,0,1),
3                                     seasonal_order= (0,1,1,12),
4                                     enforce_stationarity= False,
5                                     enforce_invertibility= False).fit()
6
7
8 predication= final_model.predict(len(df2),len(df2)+60)
9
10
11 df2.plot(legend= True, label= 'Train', figsize= (10,6))
12 predication.plot(legend= True, label= 'Prediction')
```

Out[38]: <Axes: xlabel='Order Date'>



In []: 1