

```
In [1]: 1 !pip install pmdarima==2.0.3
        2 !pip install --force-reinstall numpy==1.25.2
```

```
Collecting pmdarima==2.0.3
  Downloading pmdarima-2.0.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux_2_28_x86_64.whl.metadata (7.8 kB)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (1.4.2)
Requirement already satisfied: Cython!=0.29.18,!=0.29.31,>=0.29 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (3.0.12)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (2.0.2)
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (2.2.2)
Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (1.6.1)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (1.15.2)
Requirement already satisfied: statsmodels>=0.13.2 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (0.14.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (2.4.0)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3.11/dist-packages (from pmdarima==2.0.3) (75.2.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.19->pmdarima==2.0.3) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.19->pmdarima==2.0.3) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.19->pmdarima==2.0.3) (2025.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.22->pmdarima==2.0.3) (3.6.0)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.11/dist-packages (from statsmodels>=0.13.2->pmdarima==2.0.3) (1.0.1)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.11/dist-packages (from statsmodels>=0.13.2->pmdarima==2.0.3) (24.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas>=0.19->pmdarima==2.0.3) (1.17.0)
Downloading pmdarima-2.0.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux_2_28_x86_64.whl (1.9 MB)
----- 1.9/1.9 MB 16.8 MB/s eta 0:00:00
Installing collected packages: pmdarima
Successfully installed pmdarima-2.0.3
Collecting numpy==1.25.2
  Downloading numpy-1.25.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.6 kB)
Downloading numpy-1.25.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (18.2 MB)
----- 18.2/18.2 MB 85.0 MB/s eta 0:00:00
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 2.0.2
    Uninstalling numpy-2.0.2:
      Successfully uninstalled numpy-2.0.2
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
tensorflow 2.18.0 requires numpy<2.1.0,>=1.26.0, but you have numpy 1.25.2 which is incompatible.
thinc 8.3.6 requires numpy<3.0.0,>=2.0.0, but you have numpy 1.25.2 which is incompatible.
blosc2 3.3.2 requires numpy>=1.26, but you have numpy 1.25.2 which is incompatible.
Successfully installed numpy-1.25.2
```

Libraries

```
In [1]: 1 import pandas as pd
        2 import matplotlib.pyplot as plt
        3 from statsmodels.tsa.stattools import adfuller
        4 import numpy as np
```

Loading Data and displaying few rows data

```
In [2]: 1 df = pd.read_csv('Non seasonal nvda data.csv', parse_dates=['Date'])
        2
        3 df.head(10)
```

```
Out[2]:
```

	Date	Open	High	Low	Close	Volume
0	2025-04-07	87.46	101.75	86.62	97.64	611041250
1	2025-04-04	98.91	100.13	92.11	94.31	532273812
2	2025-04-03	103.51	105.63	101.60	101.80	338769406
3	2025-04-02	107.29	111.98	106.79	110.42	220601203
4	2025-04-01	108.52	110.20	106.47	110.15	222614000
5	2025-03-31	105.13	110.96	103.65	108.38	299212719
6	2025-03-28	111.49	112.87	109.07	109.67	229872500
7	2025-03-27	111.35	114.45	110.66	111.43	236902094
8	2025-03-26	118.73	118.84	112.71	113.76	296431719
9	2025-03-25	120.55	121.29	118.92	120.69	167447203

Data cleaning


```
In [3]: 1 # Drop empty value rows
        2 df = df.dropna()
        3
        4 # Drop duplicate rows
        5 df = df.drop_duplicates()
        6
        7 # Convert date column
        8 df['Date'] = pd.to_datetime(df['Date'])
        9
       10 # Check for Null values
       11 print(df.isnull().sum())
```

```
Date      0
Open      0
High      0
Low       0
Close     0
Volume    0
dtype: int64
```

#Data Visualization

Closing Price visual representation

```
In [4]: 1 plt.figure(figsize=(12, 6))
2 plt.plot(df['Date'], df['Close'], label='NVDA Stock Closing Price')
3 plt.title('NVDA Stock Closing Price Over Time')
4 plt.xlabel('Date')
5 plt.ylabel('Closing Price')
6 plt.legend()
7 plt.show()
```




ACF plot

```
In [5]: 1 from statsmodels.graphics.tsaplots import plot_acf
2
3 plt.figure(figsize=(12, 6))
4 acf_plot_close = plot_acf(df['Close'].dropna(), lags=48, alpha=0.05) # alpha sets confide
5
6 # Add a dotted line at the significance threshold
7 plt.axhline(y=-1.96/np.sqrt(len(df['Close'])), linestyle='--', color='gray')
8 plt.axhline(y=1.96/np.sqrt(len(df['Close'])), linestyle='--', color='gray')
9
10 plt.title('Autocorrelation Function (ACF) Plot for NVDA Close Prices with Significance Thr
11 plt.show()
```

<Figure size 1200x600 with 0 Axes>

Autocorrelation Function (ACF) Plot for NVDA Close Prices with Significance Threshold



1. Strong Positive Autocorrelation at Short Lags: What happened yesterday with the stock price really influences what happens today. They're strongly linked.
2. Gradually Decreasing Autocorrelation: As the lag increases, the autocorrelation generally decreases. This suggests that the influence of past prices on the current price weakens over time.
3. Significant Autocorrelation Beyond the Threshold: For quite a while, even as you go back several days, there's still a noticeable pattern. This suggests that the time series is not completely random and that past prices have a predictive power for future prices.

#Augmented Dickey-Fuller test (Stationary test) The ADF test is a statistical test used to determine if a time series is stationary or not.

```
In [6]: 1 result_original = adfuller(df['Close'])
2 print('ADF Statistic (Original):', result_original[0])
3 print('p-value (Original):', result_original[1])
4 print('Critical Values (Original):', result_original[4])
```

ADF Statistic (Original): -0.2988577938344866

p-value (Original): 0.9256659470708057

Critical Values (Original): {'1%': -3.4435761493506294, '5%': -2.867372960189225, '10%': -2.5698767442886696}

1. We can see that P- value is significantly greater than 0.05.
2. The ADF Statistic (-0.2988577938344866) is greater than all the critical values provided.
3. So this is non-stationary therefore we need to make it stationary by differencing.

Differencing: It helps to stabilize the mean of a time series by removing trends and seasonality.

```
In [7]: 1 #Differenced Once
        2 df['Close'] = df['Close'].diff()
```

```
In [8]: 1 df.head()
```

```
Out[8]:
```

	Date	Open	High	Low	Close	Volume
0	2025-04-07	87.46	101.75	86.62	NaN	611041250
1	2025-04-04	98.91	100.13	92.11	-3.33	532273812
2	2025-04-03	103.51	105.63	101.60	7.49	338769406
3	2025-04-02	107.29	111.98	106.79	8.62	220601203
4	2025-04-01	108.52	110.20	106.47	-0.27	222614000

After Differencing we got some NaN value in our column so we are filling it.

```
In [9]: 1 df['Close'].fillna(method='bfill', inplace=True)
```

<ipython-input-9-7170cc8430d6>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.


```
df['Close'].fillna(method='bfill', inplace=True)
<ipython-input-9-7170cc8430d6>:1: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
df['Close'].fillna(method='bfill', inplace=True)
```

```
In [10]: 1 df.head()
```

```
Out[10]:
```

	Date	Open	High	Low	Close	Volume
0	2025-04-07	87.46	101.75	86.62	-3.33	611041250
1	2025-04-04	98.91	100.13	92.11	-3.33	532273812
2	2025-04-03	103.51	105.63	101.60	7.49	338769406
3	2025-04-02	107.29	111.98	106.79	8.62	220601203
4	2025-04-01	108.52	110.20	106.47	-0.27	222614000

```
In [11]: 1 # Plotting differenced closing prices
2 plt.figure(figsize=(12, 6))
3 plt.plot(df['Date'], df['Close'], label='Differenced Closing Price')
4 plt.title('Differenced NVDA Stock Closing Price')
5 plt.xlabel('Date')
6 plt.ylabel('Differenced Closing Price')
7 plt.legend()
8 plt.show()
```




```

In [12]: 1 # ACF Plot for 'Close'
2 plt.figure(figsize=(12, 6))
3 acf_plot_close = plot_acf(df['Close'].dropna(), lags=48, alpha=0.05) # alpha sets confide
4
5 # Add a dotted line at the significance threshold
6 plt.axhline(y=-1.96/np.sqrt(len(df['Close'])), linestyle='--', color='gray')
7 plt.axhline(y=1.96/np.sqrt(len(df['Close'])), linestyle='--', color='gray')
8
9 plt.title('Autocorrelation Function (ACF) Plot for Differenced(1) Data')
10 plt.show()

```

<Figure size 1200x600 with 0 Axes>




PACF PLOT

p is the order of Autoregressive (AR). From PACF plot we can find the value of p. q is the order of Moving Average (MA). From ACF plot we can find the value of q

```
In [13]: 1 from statsmodels.graphics.tsaplots import plot_pacf
2
3 # PACF plot for 'Close'
4 plt.figure(figsize=(12, 6))
5 pacf_plot_close = plot_pacf(df['Close'].dropna(), lags=48, alpha=0.05) # Adjust 'lags' an
6
7 # Add a dotted line at the significance threshold
8 plt.axhline(y=-1.96/np.sqrt(len(df['Close'])), linestyle='--', color='gray')
9 plt.axhline(y=1.96/np.sqrt(len(df['Close'])), linestyle='--', color='gray')
10
11 plt.title('PACF Plot for NVDA Adj Close Prices after Differencing one')
12 plt.show()
```

<Figure size 1200x600 with 0 Axes>



```
In [14]: 1 from pmdarima.arima.utils import ndiffs
```



In [15]: 1 ndiffs(df['Close'], test='adf')

```

/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/deprecation.py:151: FutureWarning: 'for
ce_all_finite' was renamed to 'ensure_all_finite' in 1.6 and will be removed in 1.8.
    warnings.warn(

```

Out[15]: 0

Fitting the ARIMA Model

In [16]: 1 pip install --upgrade statsmodels

Requirement already satisfied: statsmodels in /usr/local/lib/python3.11/dist-packages (0.14.4)  
 Requirement already satisfied: numpy<3,>=1.22.3 in /usr/local/lib/python3.11/dist-packages (from statsmodels) (1.25.2)  
 Requirement already satisfied: scipy!=1.9.2,>=1.8 in /usr/local/lib/python3.11/dist-packages (from statsmodels) (1.15.2)  
 Requirement already satisfied: pandas!=2.1.0,>=1.4 in /usr/local/lib/python3.11/dist-packages (from statsmodels) (2.2.2)  
 Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.11/dist-packages (from statsmodels) (1.0.1)  
 Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.11/dist-packages (from statsmodels) (24.2)  
 Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2.9.0.post0)  
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2025.2)  
 Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2025.2)  
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas!=2.1.0,>=1.4->statsmodels) (1.17.0)

#### Model Selection

In [17]: 1 from statsmodels.tsa.arima.model import ARIMA  
 2  
 3 # ARIMA Model  
 4 model = ARIMA(df['Close'], order=(1, 1, 1))  
 5 result = model.fit()  
 6  
 7 # Print the summary  
 8 print(result.summary())

#### SARIMAX Results

```
=====
Dep. Variable:          Close    No. Observations:          501
Model:                ARIMA(1, 1, 1)    Log Likelihood          -1299.373
Date:                Mon, 05 May 2025    AIC                    2604.745
Time:                21:13:20          BIC                    2617.389
Sample:                0              HQIC                    2609.707
Covariance Type:      opg

=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
ar.L1         -0.0950     0.032     -2.922     0.003     -0.159     -0.031
ma.L1         -0.9946     0.011    -93.239     0.000     -1.016     -0.974
sigma2         10.4875     0.403     25.996     0.000      9.697     11.278
=====
Ljung-Box (L1) (Q):                0.01    Jarque-Bera (JB):                818.61
Prob(Q):                          0.90    Prob(JB):                  0.00
Heteroskedasticity (H):            0.07    Skew:                      0.65
Prob(H) (two-sided):              0.00    Kurtosis:                  9.13
=====
```

#### Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [18]: 1 from pmdarima import auto_arima
2
3 model = auto_arima(df['Close'], seasonal=False,
4                   start_p=0, start_q=0, max_p=5, max_q=5,
5                   d=1,
6                   trace=True, error_action='ignore', suppress_warnings=True, stepwise=
7 print(model.summary())
```

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.8998	0.033	-27.043	0.000	-0.965	-0.835
ar.L2	-0.6339	0.046	-13.633	0.000	-0.725	-0.543
ar.L3	-0.5334	0.046	-11.583	0.000	-0.624	-0.443
ar.L4	-0.4034	0.048	-8.346	0.000	-0.498	-0.309
ar.L5	-0.2396	0.035	-6.824	0.000	-0.308	-0.171
sigma2	12.0387	0.478	25.211	0.000	11.103	12.975


  

Ljung-Box (L1) (Q):	0.40	Jarque-Bera (JB):	613.08
Prob(Q):	0.53	Prob(JB):	0.00
Heteroskedasticity (H):	0.07	Skew:	0.55
Prob(H) (two-sided):	0.00	Kurtosis:	8.31

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

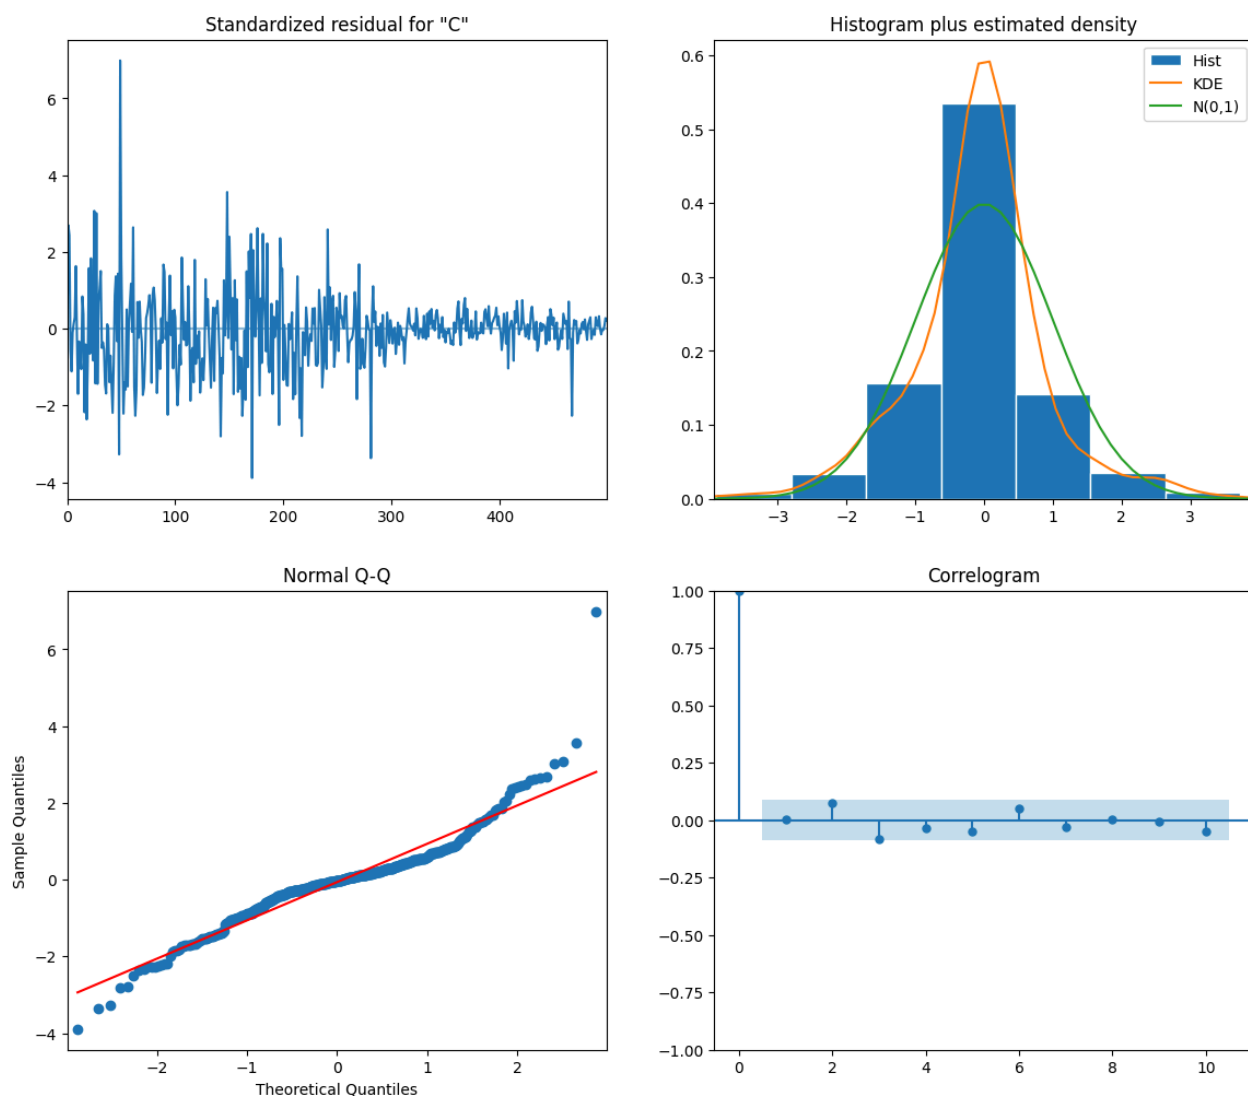
```
In [54]: 1 # Plot Residual Errors
2 residuals = pd.DataFrame(result.resid)
3
4 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 4))
5
6 ax1.plot(residuals)
7 ax1.set_title("Residuals Over Time")
8
9 ax2.hist(residuals, density=True, bins=30)
10 ax2.set_title("Distribution of Residuals")
11
12 plt.tight_layout()
13 plt.show()
14
```



```
In [50]: 1 import pandas as pd
2 from statsmodels.stats.diagnostic import acorr_ljungbox
3
4 data = df['Close']
5
6 lb_test = acorr_ljungbox(data, lags=20, return_df=True)
7
8 print(lb_test)
```

	lb_stat	lb_pvalue
1	4.512393	0.033650
2	8.500464	0.014261
3	12.273141	0.006504
4	12.342104	0.014981
5	13.367349	0.020169
6	15.331503	0.017829
7	15.981075	0.025290
8	15.986613	0.042572
9	15.987730	0.067138
10	16.630124	0.082960
11	19.345174	0.055172
12	19.684989	0.073285
13	19.897743	0.097790
14	28.117391	0.013726
15	28.265553	0.019965
16	29.993440	0.018036
17	30.074825	0.025810
18	30.401521	0.033721
19	31.323128	0.037195
20	34.233329	0.024589


```
In [22]: 1 result.plot_diagnostics(figsize=(14,12))
2 plt.show()
```



```
In [51]: 1 from scipy.stats import shapiro
2 import numpy as np
3
4 # Example dataset
5 data = residuals.copy()
6
7 # Run Shapiro-Wilk test
8 stat, p = shapiro(data)
9
10 # Print results
11 print("Shapiro-Wilk Test Statistic:", stat)
12 print("p-value:", p)
13
14 if p > 0.05:
15     print("Data appears to be normally distributed (fail to reject H0).")
16 else:
17     print("Data does not appear to be normally distributed (reject H0).")
18
```

Shapiro-Wilk Test Statistic: 0.9219184111044073  
p-value: 1.8984881534724404e-15  
Data does not appear to be normally distributed (reject H0).

```
In [34]: 1 from statsmodels.tsa.arima.model import ARIMA
2
3 # ARIMA Model
4 model = ARIMA(df['Close'], order=(1, 1, 1))
5
6 # Fit the model
7 result = model.fit()
8
9 # Get forecast
10 forecast = result.forecast(steps=20)
11
12 # Plot Actual vs Predicted
13 fig, ax = plt.subplots(figsize=(10, 6))
14 ax.plot(df['Close'], label='Actual', color='blue')
15 ax.plot(range(len(df), len(df) + len(forecast)), forecast, label='Predicted', color='red')
16 ax.legend()
17
18 plt.show()
19
20 # Print the summary
21 print(result.summary())
```



## SARIMAX Results

Dep. Variable:	Close	No. Observations:	501			
Model:	ARIMA(1, 1, 1)	Log Likelihood	-1299.373			
Date:	Mon, 05 May 2025	AIC	2604.745			
Time:	21:16:19	BIC	2617.389			
Sample:	0	HQIC	2609.707			
	- 501					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.0950	0.032	-2.922	0.003	-0.159	-0.031
ma.L1	-0.9946	0.011	-93.239	0.000	-1.016	-0.974
sigma2	10.4875	0.403	25.996	0.000	9.697	11.278
Ljung-Box (L1) (Q):	0.01	Jarque-Bera (JB):	818.61			
Prob(Q):	0.90	Prob(JB):	0.00			
Heteroskedasticity (H):	0.07	Skew:	0.65			
Prob(H) (two-sided):	0.00	Kurtosis:	9.13			

## Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [35]: 1 df2= df[['Close']].copy()
          2 df2.describe()
```

```
Out[35]:
```

	Close
count	501.000000
mean	-0.146487
std	3.247951
min	-13.290000
25%	-1.490000
50%	-0.220000
75%	1.190000
max	24.200000

## Train and Test

```
In [36]: 1 from sklearn.metrics import mean_squared_error
          2
          3 n = int(len(df2) * 0.8)
          4 train = df2[['Close'][:n]]
          5 test = df2[['Close'][n:]]
```

```
In [37]: 1 print(len(train))
          2 print(len(test))

400
101
```

```
In [38]: 1 model = ARIMA(train, order= (1,1,1))
          2 result= model.fit()
```

In [39]: 1 `print(result.summary())`

```

SARIMAX Results
=====
Dep. Variable:          Close    No. Observations:          400
Model:                ARIMA(1, 1, 1)    Log Likelihood          -1076.355
Date:                Mon, 05 May 2025    AIC                    2158.711
Time:                21:16:31    BIC                    2170.678
Sample:                0    HQIC                    2163.450
Covariance Type:          opg
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         -0.0976     0.040     -2.461     0.014     -0.175     -0.020
ma.L1         -0.9926     0.012    -84.473     0.000     -1.016     -0.970
sigma2         12.7604     0.610     20.934     0.000     11.566     13.955
=====
Ljung-Box (L1) (Q):                0.01    Jarque-Bera (JB):                394.68
Prob(Q):                0.91    Prob(JB):                0.00
Heteroskedasticity (H):                0.15    Skew:                0.66
Prob(H) (two-sided):                0.00    Kurtosis:                7.69
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```



```

In [42]: 1 # Forecast using the trained model
2 forecast_result = result.get_forecast(steps=len(test))
3 forecast_mean = forecast_result.predicted_mean
4 confidence_interval = forecast_result.conf_int()
5
6 # Plotting
7 train.plot(legend=True, label='Train', figsize=(10, 6))
8 test.plot(legend=True, label='Test')
9 forecast_mean.plot(legend=True, label='Forecast ARIMA')
10
11 # Plot confidence interval
12 plt.fill_between(confidence_interval.index, confidence_interval.iloc[:, 0], confidence_int
13
14 plt.legend()
15 plt.show()


```



```

In [43]: 1 # Forecast using the trained model
2 forecast_result = result.get_forecast(steps=len(test))
3 forecast_mean = forecast_result.predicted_mean
4 confidence_interval = forecast_result.conf_int()
5
6 # Plotting actual vs predicted values
7 plt.figure(figsize=(12, 6))
8 plt.plot(test.index, test, label='Actual')
9 plt.plot(forecast_mean.index, forecast_mean, color='red', label='Predicted')
10 plt.fill_between(confidence_interval.index, confidence_interval.iloc[:, 0], confidence_int
11 plt.title('Actual vs Predicted Values with Confidence Interval')
12 plt.xlabel('Date')
13 plt.ylabel('Close')
14 plt.legend()
15 plt.show()

```



Predict Future Data

```


In [44]: 1 final_model= ARIMA(df2, order=(1,1,1)).fit()
2
3 predication= final_model.predict(len(df2),len(df2)+60) #predicting for next 60 trading da

```

```

In [45]: 1 # Make predictions for the next 60 trading days
2 forecast_result = final_model.get_forecast(steps=60)
3 forecast_mean = forecast_result.predicted_mean
4 confidence_interval = forecast_result.conf_int()
5
6 # Plotting
7 df2.plot(legend=True, label='Train', figsize=(10, 6))
8 forecast_mean.plot(legend=True, label='Prediction')
9
10 # Plot confidence interval
11 plt.fill_between(confidence_interval.index, confidence_interval.iloc[:, 0], confidence_int
12
13 plt.legend()
14 plt.show()

```



Applying ARMA + GARCH model

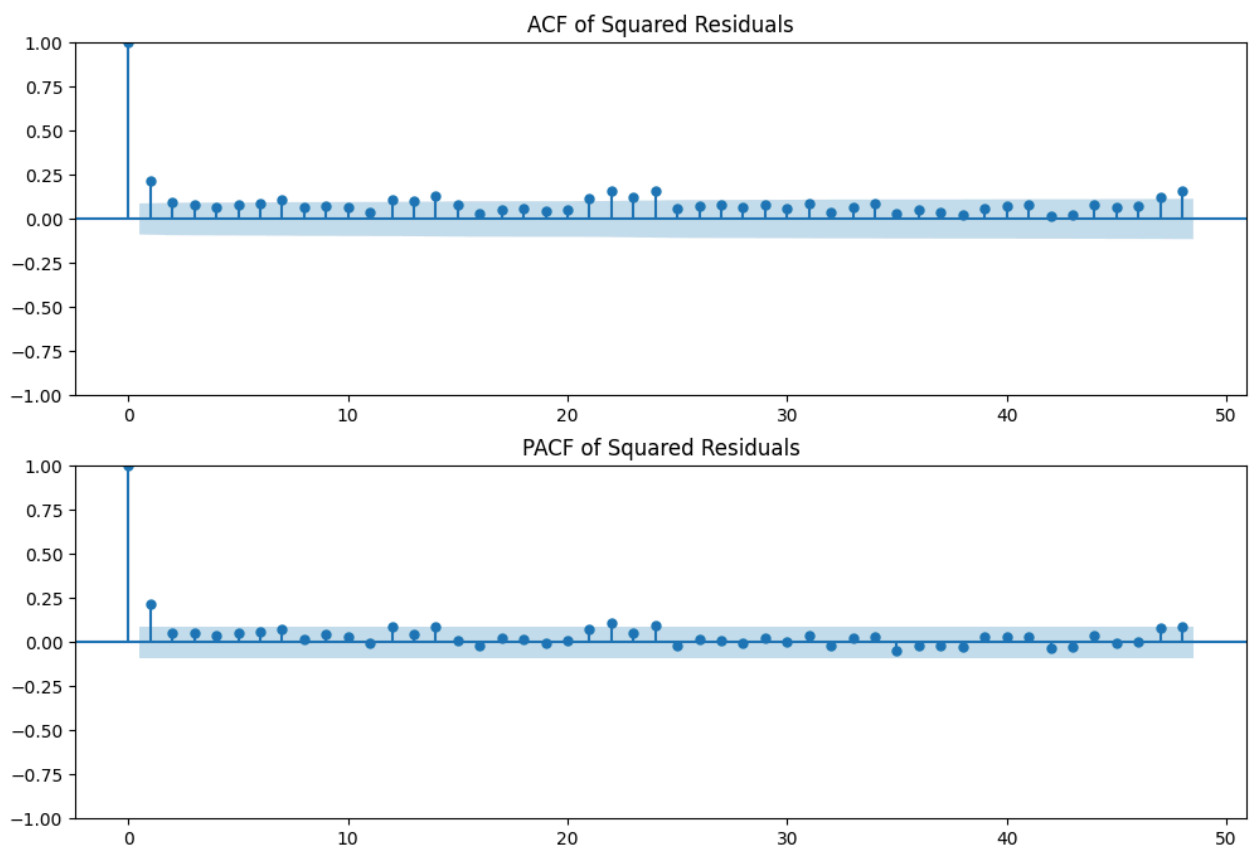
In [46]: 1 pip install arch

```
Collecting arch
  Downloading arch-7.2.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata
(13 kB)
Requirement already satisfied: numpy>=1.22.3 in /usr/local/lib/python3.11/dist-packages (from
arch) (1.25.2)
Requirement already satisfied: scipy>=1.8 in /usr/local/lib/python3.11/dist-packages (from ar
ch) (1.15.2)
Requirement already satisfied: pandas>=1.4 in /usr/local/lib/python3.11/dist-packages (from a
rch) (2.2.2)
Requirement already satisfied: statsmodels>=0.12 in /usr/local/lib/python3.11/dist-packages
(from arch) (0.14.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packa
ges (from pandas>=1.4->arch) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from
pandas>=1.4->arch) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (fro
m pandas>=1.4->arch) (2025.2)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.11/dist-packages (from
statsmodels>=0.12->arch) (1.0.1)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.11/dist-packages (fr
om statsmodels>=0.12->arch) (24.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from pyth
on-dateutil>=2.8.2->pandas>=1.4->arch) (1.17.0)
Downloading arch-7.2.0-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (985 kB)
985.3/985.3 kB 10.3 MB/s eta 0:00:00
Installing collected packages: arch
Successfully installed arch-7.2.0
```

```

In [55]: 1 # Calculate residuals
          2 residuals = final_model.resid
          3
          4 # Square the residuals
          5 squared_residuals = residuals ** 2
          6
          7 # Plot ACF and PACF of squared residuals
          8 fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 8))
          9
          10 # ACF plot
          11 plot_acf(squared_residuals, ax=ax1, lags=48)
          12 ax1.set_title('ACF of Squared Residuals')
          13
          14 # PACF plot
          15 plot_pacf(squared_residuals, ax=ax2, lags=48)
          16 ax2.set_title('PACF of Squared Residuals')
          17
          18 plt.show()

```



```
In [48]: 1 import arch
2
3 # Fit GARCH model
4 garch_model = arch.arch_model(squared_residuals, vol='Garch', p=1, q=2)
5 garch_result = garch_model.fit()
6
7 # Display the model summary
8 print(garch_result.summary())
```

```
Iteration:      1,  Func. Count:      7,  Neg. LLF: 4407.444540132237
Iteration:      2,  Func. Count:     14,  Neg. LLF: 3060.4798863927226
Iteration:      3,  Func. Count:     22,  Neg. LLF: 2651.2556633005424
Iteration:      4,  Func. Count:     29,  Neg. LLF: 2270.2972455413255
Iteration:      5,  Func. Count:     36,  Neg. LLF: 2179.308537785706
Iteration:      6,  Func. Count:     42,  Neg. LLF: 2188.481812168503
Iteration:      7,  Func. Count:     49,  Neg. LLF: 2177.3174160922003
Iteration:      8,  Func. Count:     55,  Neg. LLF: 2176.624413878339
Iteration:      9,  Func. Count:     61,  Neg. LLF: 2176.1501409323137
Iteration:     10,  Func. Count:     67,  Neg. LLF: 2174.96241701514
Iteration:     11,  Func. Count:     73,  Neg. LLF: 2372.532333221814
Iteration:     12,  Func. Count:     80,  Neg. LLF: 2394.324020554337
Iteration:     13,  Func. Count:     87,  Neg. LLF: 2278.083498148742
Iteration:     14,  Func. Count:     94,  Neg. LLF: 2419.9216926077443
Iteration:     15,  Func. Count:    101,  Neg. LLF: 2298.983680414763
Iteration:     16,  Func. Count:    108,  Neg. LLF: 2255.622829630608
Iteration:     17,  Func. Count:    115,  Neg. LLF: 2702.7731728890117
Iteration:     18,  Func. Count:    123,  Neg. LLF: 2600.018979687632
Iteration:     19,  Func. Count:    130,  Neg. LLF: 2575.215822136019
Iteration:     20,  Func. Count:    137,  Neg. LLF: 2171.0706340006586
Iteration:     21,  Func. Count:    144,  Neg. LLF: 2168.0930300650916
Iteration:     22,  Func. Count:    150,  Neg. LLF: 2168.0443351656004
Iteration:     23,  Func. Count:    156,  Neg. LLF: 2168.0350849527968
Iteration:     24,  Func. Count:    162,  Neg. LLF: 2168.032484272534
Iteration:     25,  Func. Count:    168,  Neg. LLF: 2168.029915846167
Iteration:     26,  Func. Count:    174,  Neg. LLF: 2168.0278248379836
Iteration:     27,  Func. Count:    180,  Neg. LLF: 2168.027744443383
Iteration:     28,  Func. Count:    186,  Neg. LLF: 2168.0271951683
Iteration:     29,  Func. Count:    192,  Neg. LLF: 2168.0274546278006
Optimization terminated successfully (Exit mode 0)
Current function value: 2168.027195433203
Iterations: 29
Function evaluations: 202
Gradient evaluations: 29
```

#### Constant Mean - GARCH Model Results

```
=====
Dep. Variable:      None    R-squared:      0.000
Mean Model:      Constant Mean    Adj. R-squared:      0.000
Vol Model:      GARCH    Log-Likelihood:      -2168.03
Distribution:      Normal    AIC:      4346.05
Method:      Maximum Likelihood    BIC:      4367.14
                                     No. Observations:      501
Date:      Mon, May 05 2025    Df Residuals:      500
Time:      21:16:55    Df Model:      1
                                     Mean Model
```


```
=====
              coef      std err          t      P>|t|     95.0% Conf. Int.
-----
mu              1.1214      0.829      1.353      0.176 [ -0.503,  2.746]
```

#### Volatility Model

```
=====
              coef      std err          t      P>|t|     95.0% Conf. Int.
-----
omega          5.1748     10.064      0.514      0.607 [-14.551, 24.900]
alpha[1]        0.1893    8.564e-02     2.211    2.706e-02 [2.146e-02,  0.357]
beta[1]          0.2607      0.361      0.723      0.470 [ -0.446,  0.967]
beta[2]          0.5499      0.375      1.468      0.142 [ -0.185,  1.284]
```

Covariance estimator: robust

```
In [49]: 1 # Plot results and diagnostics  
2 garch_result.plot()  
3 plt.show()
```



```
In [57]: 1 import pandas as pd
2 from arch import arch_model
3 from statsmodels.stats.diagnostic import acorr_ljungbox
4 import matplotlib.pyplot as plt
5
6 garch_model = arch_model(squared_residuals, vol='Garch', p=1, q=2)
7 garch_result = garch_model.fit()
8
9 # Display model summary
10 print(garch_result.summary())
11
12 # Get standardized residuals
13 residuals = garch_result.std_resid
14
15 # Perform Ljung-Box test on residuals
16 ljung_box = acorr_ljungbox(residuals, lags= 20, return_df=True)
17
18 print("\nLjung-Box Test Results:")
19 print(ljung_box)
20
```



```

Iteration:      1,  Func. Count:      7,  Neg. LLF: 4407.444540132237
Iteration:      2,  Func. Count:     14,  Neg. LLF: 3060.4798863927226
Iteration:      3,  Func. Count:     22,  Neg. LLF: 2651.2556633005424
Iteration:      4,  Func. Count:     29,  Neg. LLF: 2270.2972455413255
Iteration:      5,  Func. Count:     36,  Neg. LLF: 2179.308537785706
Iteration:      6,  Func. Count:     42,  Neg. LLF: 2188.481812168503
Iteration:      7,  Func. Count:     49,  Neg. LLF: 2177.3174160922003
Iteration:      8,  Func. Count:     55,  Neg. LLF: 2176.624413878339
Iteration:      9,  Func. Count:     61,  Neg. LLF: 2176.1501409323137
Iteration:     10,  Func. Count:     67,  Neg. LLF: 2174.96241701514
Iteration:     11,  Func. Count:     73,  Neg. LLF: 2372.532333221814
Iteration:     12,  Func. Count:     80,  Neg. LLF: 2394.324020554337
Iteration:     13,  Func. Count:     87,  Neg. LLF: 2278.083498148742
Iteration:     14,  Func. Count:     94,  Neg. LLF: 2419.9216926077443
Iteration:     15,  Func. Count:    101,  Neg. LLF: 2298.983680414763
Iteration:     16,  Func. Count:    108,  Neg. LLF: 2255.622829630608
Iteration:     17,  Func. Count:    115,  Neg. LLF: 2702.7731728890117
Iteration:     18,  Func. Count:    123,  Neg. LLF: 2600.018979687632
Iteration:     19,  Func. Count:    130,  Neg. LLF: 2575.215822136019
Iteration:     20,  Func. Count:    137,  Neg. LLF: 2171.0706340006586
Iteration:     21,  Func. Count:    144,  Neg. LLF: 2168.0930300650916
Iteration:     22,  Func. Count:    150,  Neg. LLF: 2168.0443351656004
Iteration:     23,  Func. Count:    156,  Neg. LLF: 2168.0350849527968
Iteration:     24,  Func. Count:    162,  Neg. LLF: 2168.032484272534
Iteration:     25,  Func. Count:    168,  Neg. LLF: 2168.029915846167
Iteration:     26,  Func. Count:    174,  Neg. LLF: 2168.0278248379836
Iteration:     27,  Func. Count:    180,  Neg. LLF: 2168.027744443383
Iteration:     28,  Func. Count:    186,  Neg. LLF: 2168.0271951683
Iteration:     29,  Func. Count:    192,  Neg. LLF: 2168.0274546278006

```

Optimization terminated successfully (Exit mode 0)

Current function value: 2168.027195433203

Iterations: 29

Function evaluations: 202

Gradient evaluations: 29

Constant Mean – GARCH Model Results

```

=====
Dep. Variable:      None      R-squared:      0.000
Mean Model:         Constant Mean  Adj. R-squared:      0.000
Vol Model:          GARCH      Log-Likelihood:    -2168.03
Distribution:        Normal    AIC:              4346.05
Method:             Maximum Likelihood  BIC:              4367.14
                                     No. Observations:    501
Date:               Mon, May 05 2025  Df Residuals:      500
Time:               21:52:47          Df Model:         1
                                     Mean Model
=====

```

```

=====
              coef      std err          t      P>|t|     95.0% Conf. Int.
-----
mu           1.1214      0.829        1.353     0.176 [ -0.503,  2.746]
=====
              Volatility Model
=====

```

```

=====
              coef      std err          t      P>|t|     95.0% Conf. Int.
-----
omega        5.1748     10.064         0.514     0.607 [-14.551, 24.900]
alpha[1]     0.1893     8.564e-02        2.211  2.706e-02 [2.146e-02,  0.357]
beta[1]      0.2607      0.361         0.723     0.470 [ -0.446,  0.967]
beta[2]      0.5499      0.375         1.468     0.142 [ -0.185,  1.284]
=====

```

Covariance estimator: robust

Ljung-Box Test Results:

```

lb_stat  lb_pvalue
1      3.001522  0.083186
2      6.009235  0.049558
3      6.239033  0.100543
4      6.693593  0.152994
5      6.874028  0.230179
6     15.229046  0.018548
7     15.680577  0.028200
8     16.352810  0.037600

```

9	17.499147	0.041450
10	17.606637	0.061973
11	18.273491	0.075448
12	19.236631	0.082976
13	24.962605	0.023346
14	25.476277	0.030146
15	26.309785	0.034890
16	26.317864	0.049717
17	26.503731	0.065759
18	26.551016	0.087810
19	27.164798	0.100873
20	27.712404	0.116384

In [74]:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 lags = ljung_box.index.to_numpy()
5 p_values = ljung_box['lb_pvalue'].to_numpy()
6
7 plt.plot(lags, p_values, marker='o', color='red')
8 plt.xlabel('Lag')
9 plt.ylabel('p-value')
10 plt.title('Ljung-Box Test Results')
11 plt.axhline(y=0.05, color='blue', linestyle='--', label='Significance Level (0.05)')
12 plt.legend()
13 plt.grid(True)
14 plt.show()
15
```

